

# SQL vs NoSQL

Νίκος Βορνιωτάκης

June 9, 2016

# Περιεχόμενα

## Εισαγωγή

Έννοιες

Relational

## Τύποι NoSQL βάσεων

Document

Keystore

Column

Graph

## Use Case

Wikipedia

Twitter

Others

## Conclusion

Conclusions

# Layout

## Εισαγωγή

Έννοιες

Relational

## Τύποι NoSQL βάσεων

Document

Keystore

Column

Graph

## Use Case

Wikipedia

Twitter

Others

## Conclusion

Conclusions

## NoSQL Databases

APACHE  
**HBASE**



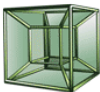
*Cassandra*



**CouchDB**  
relax



**riak**



**mongoDB**

**HYPERTABLE** INC



**Neo4j**



**redis**

# NoSQL

## Definition

- Originally referring to "non SQL" or "non relational"

# NoSQL

## Definition

- Originally referring to "non SQL" or "non relational"
- Sometimes called "Not only SQL" to emphasize that they may support SQL-like query languages

# NoSQL

## Definition

- Originally referring to "non SQL" or "non relational"
- Sometimes called "Not only SQL" to emphasize that they may support SQL-like query languages

## Characteristics

- Distributed (designed to scale)

# NoSQL

## Definition

- Originally referring to "non SQL" or "non relational"
- Sometimes called "Not only SQL" to emphasize that they may support SQL-like query languages

## Characteristics

- Distributed (designed to scale)
- Web friendly (RESTfull, JSON aware etc)



# NoSQL

## Definition

- Originally referring to "non SQL" or "non relational"
- Sometimes called "Not only SQL" to emphasize that they may support SQL-like query languages

## Characteristics

- Distributed (designed to scale)
- Web friendly (RESTfull, JSON aware etc)
- No strict schema

# NoSQL

## Definition

- Originally referring to "non SQL" or "non relational"
- Sometimes called "Not only SQL" to emphasize that they may support SQL-like query languages

## Characteristics

- Distributed (designed to scale)
- Web friendly (RESTfull, JSON aware etc)
- No strict schema
- Can support BIGDATA

# NoSQL

## Definition

- Originally referring to "non SQL" or "non relational"
- Sometimes called "Not only SQL" to emphasize that they may support SQL-like query languages

## Characteristics

- Distributed (designed to scale)
- Web friendly (RESTfull, JSON aware etc)
- No strict schema
- Can support BIGDATA
- Performance gains usually when many nodes are used

# NoSQL

## Definition

- Originally referring to "non SQL" or "non relational"
- Sometimes called "Not only SQL" to emphasize that they may support SQL-like query languages

## Characteristics

- Distributed (designed to scale)
- Web friendly (RESTfull, JSON aware etc)
- No strict schema
- Can support BIGDATA
- Performance gains usually when many nodes are used
- Suitable for commodity hardware

## CAP theorem

### Theorem

*It is impossible for a distributed computer system to simultaneously provide all three of the following guarantees:*

## CAP theorem

### Theorem

*It is impossible for a distributed computer system to simultaneously provide all three of the following guarantees:*

- *Consistency (all nodes see the same data at the same time)*

## CAP theorem

### Theorem

*It is impossible for a distributed computer system to simultaneously provide all three of the following guarantees:*

- *Consistency (all nodes see the same data at the same time)*
- *Availability (a guarantee that every request receives a response about whether it succeeded or failed)*

## CAP theorem

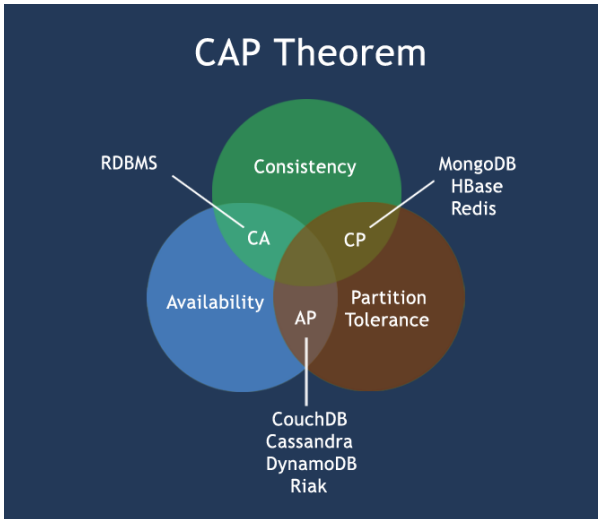
### Theorem

*It is impossible for a distributed computer system to simultaneously provide all three of the following guarantees:*

- *Consistency (all nodes see the same data at the same time)*
- *Availability (a guarantee that every request receives a response about whether it succeeded or failed)*
- *Partition tolerance (the system continues to operate despite arbitrary partitioning due to network failures)*



## CAP theorem



# ACID

Transactional guarantees

# ACID

## Transactional guarantees

**Atomicity** each transaction is "all or nothing": if one part of the transaction fails, then the entire transaction fails, and the database state is left unchanged

# ACID

## Transactional guarantees

**Atomicity** each transaction is "all or nothing": if one part of the transaction fails, then the entire transaction fails, and the database state is left unchanged

**Consistency** any transaction will bring the database from one valid state to another

# ACID

## Transactional guarantees

**Atomicity** each transaction is "all or nothing": if one part of the transaction fails, then the entire transaction fails, and the database state is left unchanged

**Consistency** any transaction will bring the database from one valid state to another

**Isolation** concurrent execution of transactions results in a system state that would be obtained if transactions were executed serially

# ACID

## Transactional guarantees

**Atomicity** each transaction is "all or nothing": if one part of the transaction fails, then the entire transaction fails, and the database state is left unchanged

**Consistency** any transaction will bring the database from one valid state to another

**Isolation** concurrent execution of transactions results in a system state that would be obtained if transactions were executed serially

**Durability** once a transaction has been committed, it will remain so, even in the event of power loss, crashes, or errors

# ACID

## Transactional guarantees

**Atomicity** each transaction is "all or nothing": if one part of the transaction fails, then the entire transaction fails, and the database state is left unchanged

**Consistency** any transaction will bring the database from one valid state to another

**Isolation** concurrent execution of transactions results in a system state that would be obtained if transactions were executed serially

**Durability** once a transaction has been committed, it will remain so, even in the event of power loss, crashes, or errors

# ACID

Transactional guarantees

**RDBMSs are Transactional**

RDBMSs provide ACID guarantees.



# ACID

## Transactional guarantees

### RDBMSs are Transactional

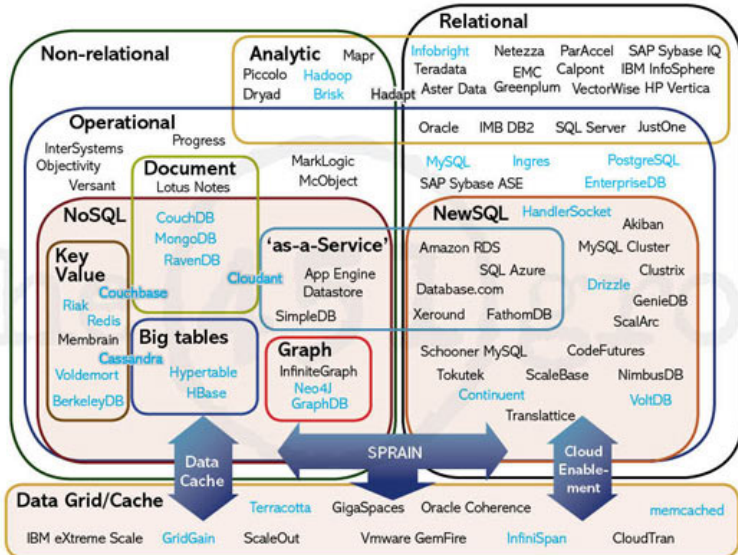
RDBMSs provide ACID guarantees.

### Most NoSQL are **not** Transactional

Most NoSQL do not provide ACID.

- On a single node ACID is "simple"
- On distributed systems to provide ACID guarantees is **really** hard

# Database Classification



Εισαγωγή	Τύποι NoSQL βάσεων	Use Case	Conclusion
oooooooo●	oooooooooooo	oooo	ooooo
oooooooooooo	ooo	oo	
	ooo	oo	
	ooooooooo		

# NewSQL

## Characteristics

- Relational

# NewSQL

## Characteristics

- Relational
- Distributed (designed to scale)

# NewSQL

## Characteristics

- Relational
- Distributed (designed to scale)
- ACID

# NewSQL

## Characteristics

- Relational
- Distributed (designed to scale)
- ACID
- Circumvents CAP by using distributed subsystems

# NewSQL

## Characteristics

- Relational
- Distributed (designed to scale)
- ACID
- Circumvents CAP by using distributed subsystems
- Relies on extremely complex distributed algorithms to maintain consistency. Heavy use of hardware-assisted time synchronization using GPS clocks and atomic clocks to ensure global consistency.

# Layout

## Εισαγωγή

Έννοιες

Relational

## Τύποι NoSQL βάσεων

Document

Keystore

Column

Graph

## Use Case

Wikipedia

Twitter

Others

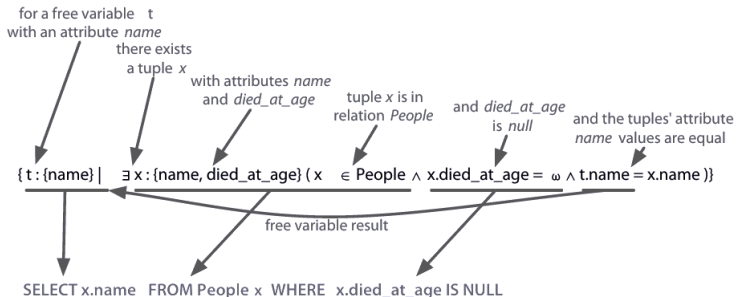
## Conclusion

Conclusions



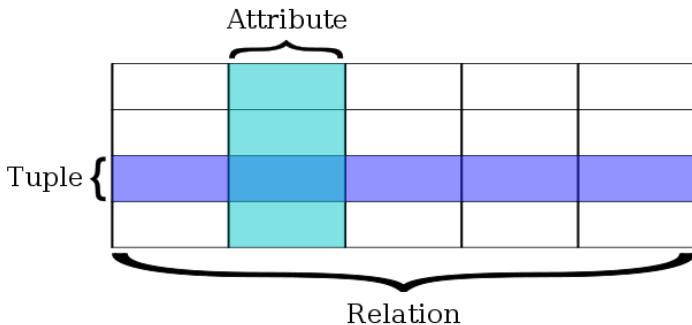
## Relational model

Relational model comes straight from relational algebra



## Relational model

Relation (relational algebra) = Table (SQL)



# Relational model

## Structured Query Language (SQL)

### A language to make queries.

The current ISO SQL standard doesn't mention the relational model or use relational terms or concepts. However, it is possible to create a database conforming to the relational model using SQL if one does not use certain SQL features.

# Relational model

## Structured Query Language (SQL)

### A language to make queries.

The current ISO SQL standard doesn't mention the relational model or use relational terms or concepts. However, it is possible to create a database conforming to the relational model using SQL if one does not use certain SQL features.

```
SELECT Book.title AS Title,  
        count(*) AS Authors  
FROM   Book  
JOIN   Book_author  
  ON    Book.isbn = Book_author.isbn  
GROUP BY Book.title;
```

# Relational model

## Structured Query Language (SQL)

### A language to make queries.

The current ISO SQL standard doesn't mention the relational model or use relational terms or concepts. However, it is possible to create a database conforming to the relational model using SQL if one does not use certain SQL features.

```
SELECT Book.title AS Title,  
        count(*) AS Authors  
FROM   Book  
JOIN   Book_author  
        ON   Book.isbn = Book_author.isbn  
GROUP BY Book.title;
```

### output

Title	Authors
SQL Examples and Guide	4
The Joy of SQL	1
An Introduction to SQL	2
Pitfalls of SQL	1

## Relational model

### Relational Model

Activity Code	Activity Name
23	Patching
24	Overlay
25	Crack Sealing

Key = 24

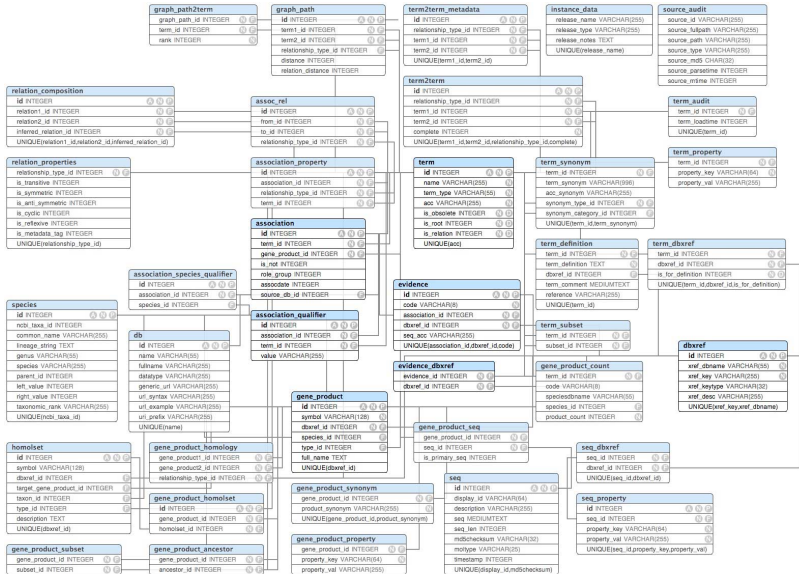
Activity Code	Date	Route No.
24	01/12/01	I-95
24	02/08/01	I-66

Date	Activity Code	Route No.
01/12/01	24	I-95
01/15/01	23	I-495
02/08/01	24	I-66

Tables are split to smallest possible not containing repeated data (normalisation) and relationships are defined among tables to create the relational **schema...**

# Relational model

## Relational schema



## Relational model

### Schema drawbacks

- Relational model is not easy to grasp

$$(R \bowtie S) \cup ((R - \pi_{r_1, r_2, \dots, r_n}(R \bowtie S)) \times (\omega, \dots \omega))$$



## Relational model

### Schema drawbacks

- Relational model is not easy to grasp  

$$(R \bowtie S) \cup ((R - \pi_{r_1, r_2, \dots, r_n}(R \bowtie S)) \times (\omega, \dots \omega))$$
- It may be hard to transform application data model to a relational schema

## Relational model

### Schema drawbacks

- Relational model is not easy to grasp  

$$(R \bowtie S) \cup ((R - \pi_{r_1, r_2, \dots, r_n}(R \bowtie S)) \times (\omega, \dots \omega))$$
- It may be hard to transform application data model to a relational schema
- Schema changes are hard to perform

## Relational model

### Schema drawbacks

- Relational model is not easy to grasp  

$$(R \bowtie S) \cup ((R - \pi_{r_1, r_2, \dots, r_n}(R \bowtie S)) \times (\omega, \dots \omega))$$
- It may be hard to transform application data model to a relational schema
- Schema changes are hard to perform
- It can be difficult (or even impossible) to model some relations to relational model (eg tree structures, transitive closure etc.)

# RDBMS

## Most popular

Oracle, MySQL, Microsoft SQL Server, PostgreSQL and IBM DB2

# RDBMS

## Most popular

Oracle, MySQL, Microsoft SQL Server, PostgreSQL and IBM DB2

- Great query flexibility

# RDBMS

## Most popular

Oracle, MySQL, Microsoft SQL Server, PostgreSQL and IBM DB2

- Great query flexibility
- Data consistency and durability ACID

# RDBMS

## Most popular

Oracle, MySQL, Microsoft SQL Server, PostgreSQL and IBM DB2

- Great query flexibility
- Data consistency and durability ACID
- Tables modeled using relational algebra

# RDBMS

## Most popular

Oracle, MySQL, Microsoft SQL Server, PostgreSQL and IBM DB2

- Great query flexibility
- Data consistency and durability ACID
- Tables modeled using relational algebra
- Data is normalised (space efficient)



## RDBMS drawbacks

## RDBMS drawbacks

- Not good at partitioning

## RDBMS drawbacks

- Not good at partitioning
- Not designed for horizontal scaling

## RDBMS drawbacks

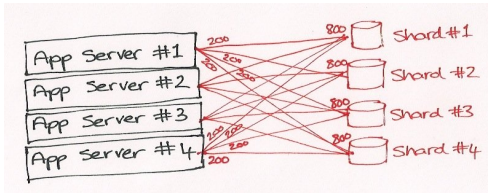
- Not good at partitioning
- Not designed for horizontal scaling
- Your data may be hard to model in strict relational model, or just not worth it

## RDBMS drawbacks

- Not good at partitioning
- Not designed for horizontal scaling
- Your data may be hard to model in strict relational model, or just not worth it
- You need to store large data blobs

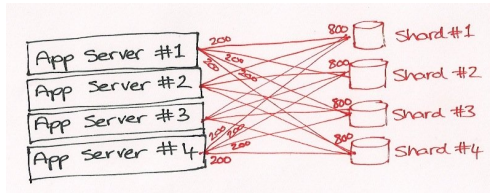
# RDBMS scaling

Sharding for partitioning



# RDBMS scaling

Sharding for partitioning

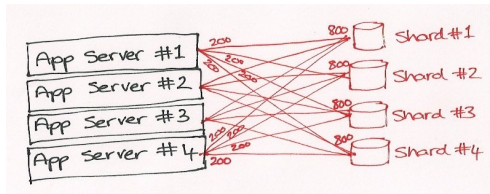


Large Datasets can be sharded for scaling

- Poor man's partitioning

# RDBMS scaling

## Sharding for partitioning



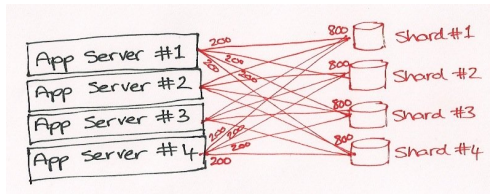
## Large Datasets can be sharded for scaling

- Poor man's partitioning
- Relies on application code (bad™)



# RDBMS scaling

## Sharding for partitioning

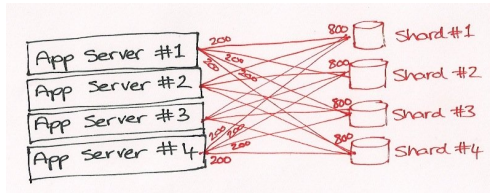


## Large Datasets can be sharded for scaling

- Poor man's partitioning
- Relies on application code (bad™)
- May work for small number of tables

# RDBMS scaling

## Sharding for partitioning

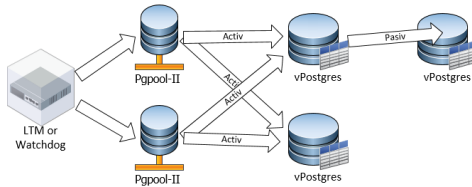


## Large Datasets can be sharded for scaling

- Poor man's partitioning
- Relies on application code (bad™)
- May work for small number of tables
- Not easy to add nodes

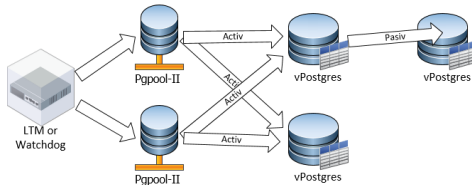
# RDBMS scaling

Replication for availability



# RDBMS scaling

Replication for availability

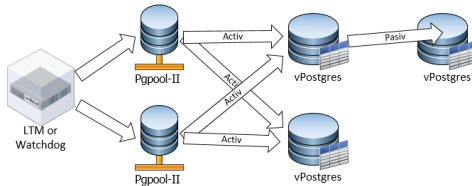


Replication can be used to improve small datasets  
READ/WRITE performance

- Usually requires 3rd party solutions

# RDBMS scaling

Replication for availability

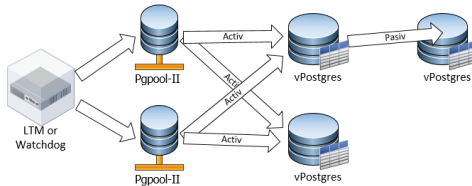


Replication can be used to improve small datasets  
READ/WRITE performance

- Usually requires 3rd party solutions
- Complicated setup

# RDBMS scaling

## Replication for availability



Replication can be used to improve small datasets  
READ/WRITE performance

- Usually requires 3rd party solutions
- Complicated setup
- WRITE has problems (blocking)

# RDBMS scaling

Sharding and replication?

Sharding and replication for big production systems has many problems!

- Partitioning is hard, solutions rely on application code (BAD™)

# RDBMS scaling

Sharding and replication?

Sharding and replication for big production systems has many problems!

- Partitioning is hard, solutions rely on application code (BAD™)
- IO **will** be a bottleneck no matter what. Multiple caching layers need to be designed.



# RDBMS scaling

Sharding and replication?

Sharding and replication for big production systems has many problems!

- Partitioning is hard, solutions rely on application code (BAD™)
- IO **will** be a bottleneck no matter what. Multiple caching layers need to be designed.
- Architecture becomes really complicated. Hard to manage.

# RDBMS scaling

Sharding and replication?

Sharding and replication for big production systems has many problems!

- Partitioning is hard, solutions rely on application code (BAD™)
- IO **will** be a bottleneck no matter what. Multiple caching layers need to be designed.
- Architecture becomes really complicated. Hard to manage.
- Not part of DBMS. Need to rely on heterogeneous technologies.

# RDBMS scaling

Sharding and replication?

Sharding and replication for big production systems has many problems!

- Partitioning is hard, solutions rely on application code (BAD™)
- IO **will** be a bottleneck no matter what. Multiple caching layers need to be designed.
- Architecture becomes really complicated. Hard to manage.
- Not part of DBMS. Need to rely on heterogeneous technologies.

# RDBMS scaling

Sharding and replication?

Sharding and replication for big production systems has many problems!

- Partitioning is hard, solutions rely on application code (BAD™)
- IO **will** be a bottleneck no matter what. Multiple caching layers need to be designed.
- Architecture becomes really complicated. Hard to manage.
- Not part of DBMS. Need to rely on heterogeneous technologies.

NoSQL can help

# Layout

## Εισαγωγή

Έννοιες

Relational

## Τύποι NoSQL βάσεων

Document

Keystore

Column

Graph

## Use Case

Wikipedia

Twitter

Others

## Conclusion

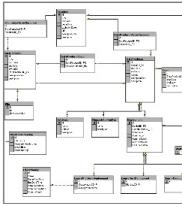
Conclusions

# Document

Most popular

MongoDB, CouchDB

RDBMS



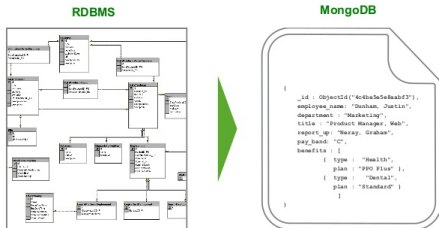
MongoDB



# Document

Most popular

MongoDB, CouchDB



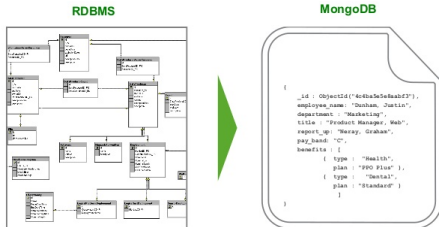
## Θετικά

- Horizontal Scaling and performance
- Hierarchical data
- Varied data (no schema)
- Web friendly (RESTfull, Javascript)
- Suitable for large data blobs

# Document

Most popular

MongoDB, CouchDB



## Θετικά

- Horizontal Scaling and performance
- Hierarchical data
- Varied data (no schema)
- Web friendly (RESTfull, Javascript)
- Suitable for large data blobs

## Αρνητικά

- No Transactions
- Eventual consistency
- Varied data (no schema)
- Denormalised data
- Needs careful document modeling (for big scale)



## Mongo Document Example

```
{
  "_id" : ObjectId("54c955492b7c8eb21818bd09"),
  "address" : {
    "street" : "2 Avenue",
    "zipcode" : "10075",
    "building" : "1480",
    "coord" : [ -73.9557413, 40.7720266 ]
  },
  "borough" : "Manhattan",
  "cuisine" : "Italian",
  "grades" : [
    {
      "date" : ISODate("2014-10-01T00:00:00Z"),
      "grade" : "A",
      "score" : 11
    },
    {
      "date" : ISODate("2014-01-16T00:00:00Z"),
      "grade" : "B",
      "score" : 17
    }
  ],
  "name" : "Vella",
  "restaurant_id" : "41704620"
}
```

## Mongo INSERT

```
> db.towns.insert({  
  name: "New York",  
  population: 22200000,  
  last_census: ISODate("2009-07-31"),  
  famous_for: [ "statue of liberty", "food" ],  
  mayor : {  
    name : "Michael Bloomberg",  
    party : "I"  
  }  
})
```

## Mongo SELECT

```
> db.towns.find()
{
  "_id" : ObjectId("4d0ad975bb30773266f39fe3"),
  "name" : "New York",
  "population": 22200000,
  "last_census": "Fri Jul 31 2009 00:00:00 GMT-0700 (PDT)",
  "famous_for" : [ "statue of liberty", "food" ],
  "mayor" : { "name" : "Michael Bloomberg", "party" : "I" }
}
```

## CouchDB SELECT \*

```
# curl http://localhost:5984/music/  
  
{  
  "db_name": "music",  
  "doc_count": 1,  
  "doc_del_count": 0,  
  "update_seq": 4,  
  "purge_seq": 0,  
  "compact_running": false,  
  "disk_size": 16473,  
  "instance_start_time": "1326845777510067",  
  "disk_format_version": 5,  
  "committed_update_seq": 4  
}
```

## CouchDB SELECT

```
# curl http://localhost:5984/music/74c7a8d2a8548c8b97da748f43000ac4",
{
  "_id": "74c7a8d2a8548c8b97da748f43000ac4",
  "_rev": "4-93a101178ba65f61ed39e60d70c9fd97",
  "name": "The Beatles",
  "albums": [
    {
      "title": "Help!",
      "year": 1965
    }, {
      "title": "Sgt. Pepper's Lonely Hearts Club Band",
      "year": 1967
    }, {
      "title": "Abbey Road",
      "year": 1969
    }
  ]
}
```

## CouchDB INSERT

```
# curl -i -X POST "http://localhost:5984/music/" \
-H "Content-Type: application/json" \
-d '{ "name": "Wings" }'
```

HTTP/1.1 201 Created

Server: CouchDB/1.1.1 (Erlang OTP/R14B03)

Location: <http://localhost:5984/music/74c7a8d2a8548c8b97da748f43000f1b>

Date: Wed, 18 Jan 2012 00:37:51 GMT

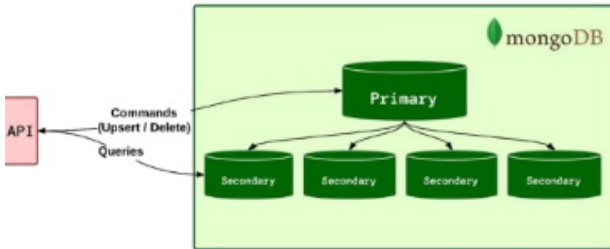
Content-Type: text/plain; charset=utf-8

Content-Length: 95

Cache-Control: must-revalidate

```
{
  "ok": true,
  "id": "74c7a8d2a8548c8b97da748f43000f1b",
  "rev": "1-2fe1dd1911153eb9df8460747dfe75a0"
}
```

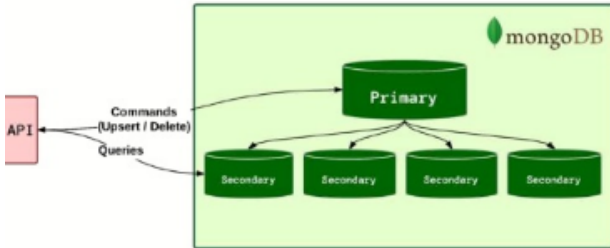
# Mongo scaling with replica sets



## Characteristics

- Built in the core – API agnostic

# Mongo scaling with replica sets

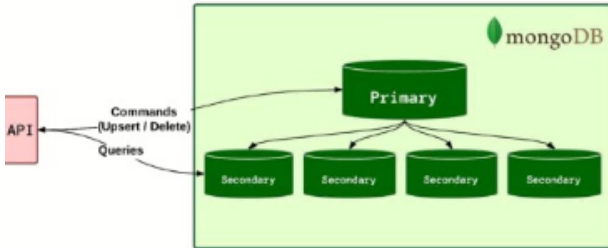


## Characteristics

- Built in the core – API agnostic
- Redundancy — auto-failover



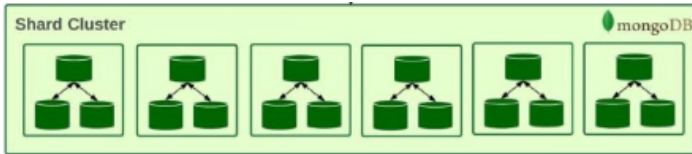
# Mongo scaling with replica sets



## Characteristics

- Built in the core – API agnostic
- Redundancy — auto-failover
- Easy to setup

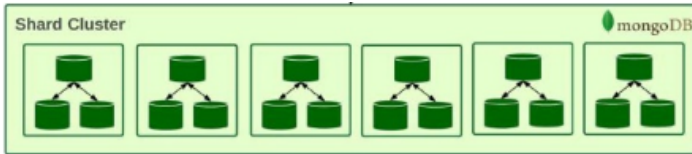
# Mongo scaling with sharding



## Characteristics

- Each piece of data is exclusively controlled by a single node (shard)

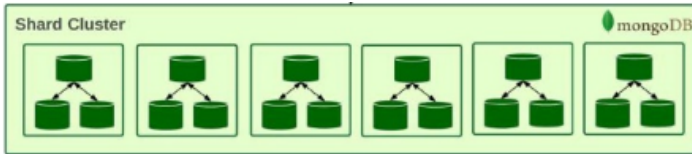
# Mongo scaling with sharding



## Characteristics

- Each piece of data is exclusively controlled by a single node (shard)
- Each node (shard) has exclusive control over a well defined subset of the data

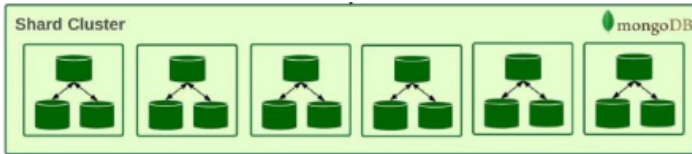
# Mongo scaling with sharding



## Characteristics

- Each piece of data is exclusively controlled by a single node (shard)
- Each node (shard) has exclusive control over a well defined subset of the data
- As system load changes, assignment of data to shards is **rebalanced** automatically

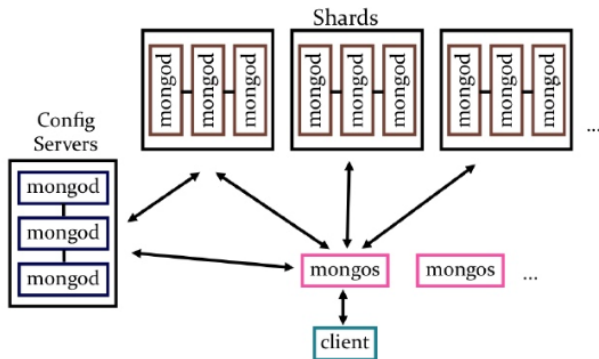
# Mongo scaling with sharding



## Characteristics

- Each piece of data is exclusively controlled by a single node (shard)
- Each node (shard) has exclusive control over a well defined subset of the data
- As system load changes, assignment of data to shards is **rebalanced** automatically
- Can be combined with replica set for per-shard redundancy

# Mongo system architecture



Combining sharding and replication for complicate architectures.  
All nodes are vanilla Mongos.

# Mongo scaling caveats

## scaling caveats

- Relatively easy to setup and only using one type of software.  
No need for third party solutions.

## Mongo scaling caveats

### scaling caveats

- Relatively easy to setup and only using one type of software.  
No need for third party solutions.
- While schema-free design is flexible and some times can model the application data model better than the relational model, at scale you need to carefully model your Documents to not end up with great performance loss.



## Mongo scaling caveats

### scaling caveats

- Relatively easy to setup and only using one type of software.  
No need for third party solutions.
- While schema-free design is flexible and some times can model the application data model better than the relational model, at scale you need to carefully model your Documents to not end up with great performance loss.
- If sharding is used, a well thought sharding key must be used from the beginning since it is very difficult to change.

## Mongo scaling caveats

### scaling caveats

- Relatively easy to setup and only using one type of software.  
No need for third party solutions.
- While schema-free design is flexible and some times can model the application data model better than the relational model, at scale you need to carefully model your Documents to not end up with great performance loss.
- If sharding is used, a well thought sharding key must be used from the beginning since it is very difficult to change.
- For production systems, the minimum nodes required for a scalable architecture is 9. Are your database needs that large?

# Layout

## Εισαγωγή

Έννοιες

Relational

## Τύποι NoSQL βάσεων

Document

**Keystore**

Column

Graph

## Use Case

Wikipedia

Twitter

Others

## Conclusion

Conclusions

# Keystore

Most popular

Riak, Redis, memcached

Key	Value
K1	AAA,BBB,CCC
K2	AAA,BBB
K3	AAA,DDD
K4	AAA,2,01/01/2015
K5	3,ZZZ,5623

# Keystore

Most popular

Riak, Redis, memcached

Χαρακτηριστικά

- No Transactions

Key	Value
K1	AAA,BBB,CCC
K2	AAA,BBB
K3	AAA,DDD
K4	AAA,2,01/01/2015
K5	3,ZZZ,5623

# Keystore

## Most popular

Riak, Redis, memcached

## Χαρακτηριστικά

- No Transactions
- No SQL

Key	Value
K1	AAA,BBB,CCC
K2	AAA,BBB
K3	AAA,DDD
K4	AAA,2,01/01/2015
K5	3,ZZZ,5623

# Keystore

## Most popular

Riak, Redis, memcached

## Χαρακτηριστικά

- No Transactions
- No SQL
- No Schema

Key	Value
K1	AAA,BBB,CCC
K2	AAA,BBB
K3	AAA,DDD
K4	AAA,2,01/01/2015
K5	3,ZZZ,5623

# Keystore

## Most popular

Riak, Redis, memcached

## Χαρακτηριστικά

- No Transactions
- No SQL
- No Schema

Key	Value
K1	AAA,BBB,CCC
K2	AAA,BBB
K3	AAA,DDD
K4	AAA,2,01/01/2015
K5	3,ZZZ,5623



# Keystore

## Most popular

Riak, Redis, memcached

## Χαρακτηριστικά

- No Transactions
- No SQL
- No Schema

Key	Value
K1	AAA,BBB,CCC
K2	AAA,BBB
K3	AAA,DDD
K4	AAA,2,01/01/2015
K5	3,ZZZ,5623

## Θετικά

- Horizontal Scaling

# Keystore

## Most popular

Riak, Redis, memcached

## Χαρακτηριστικά

- No Transactions
- No SQL
- No Schema

Key	Value
K1	AAA,BBB,CCC
K2	AAA,BBB
K3	AAA,DDD
K4	AAA,2,01/01/2015
K5	3,ZZZ,5623

## Θετικά

- Horizontal Scaling
- High Availability

# Keystore

## Most popular

Riak, Redis, memcached

## Χαρακτηριστικά

- No Transactions
- No SQL
- No Schema

Key	Value
K1	AAA,BBB,CCC
K2	AAA,BBB
K3	AAA,DDD
K4	AAA,2,01/01/2015
K5	3,ZZZ,5623

## Θετικά

- Horizontal Scaling
- High Availability
- Extremely fast by design

# Keystore

## Most popular

Riak, Redis, memcached

## Χαρακτηριστικά

- No Transactions
- No SQL
- No Schema

## Θετικά

- Horizontal Scaling
- High Availability
- Extremely fast by design
- Usually in memory

Key	Value
K1	AAA,BBB,CCC
K2	AAA,BBB
K3	AAA,DDD
K4	AAA,2,01/01/2015
K5	3,ZZZ,5623

# Keystore

## Most popular

Riak, Redis, memcached

## Χαρακτηριστικά

- No Transactions
- No SQL
- No Schema

## Θετικά

- Horizontal Scaling
- High Availability
- Extremely fast by design
- Usually in memory

Key	Value
K1	AAA,BBB,CCC
K2	AAA,BBB
K3	AAA,DDD
K4	AAA,2,01/01/2015
K5	3,ZZZ,5623

# Keystore

## Most popular

Riak, Redis, memcached

## Χαρακτηριστικά

- No Transactions
- No SQL
- No Schema

## Αρνητικά

- Lack of schema

Key	Value
K1	AAA,BBB,CCC
K2	AAA,BBB
K3	AAA,DDD
K4	AAA,2,01/01/2015
K5	3,ZZZ,5623

## Θετικά

- Horizontal Scaling
- High Availability
- Extremely fast by design
- Usually in memory

# Keystore

## Most popular

Riak, Redis, memcached

## Χαρακτηριστικά

- No Transactions
- No SQL
- No Schema

## Αρνητικά

- Lack of schema
- Lack of ACID

Key	Value
K1	AAA,BBB,CCC
K2	AAA,BBB
K3	AAA,DDD
K4	AAA,2,01/01/2015
K5	3,ZZZ,5623

## Θετικά

- Horizontal Scaling
- High Availability
- Extremely fast by design
- Usually in memory

# Keystore

## Most popular

Riak, Redis, memcached

## Χαρακτηριστικά

- No Transactions
- No SQL
- No Schema

## Αρνητικά

- Lack of schema
- Lack of ACID
- Lack of robust querying, only basic CRUD

## Θετικά

- Horizontal Scaling
- High Availability
- Extremely fast by design
- Usually in memory

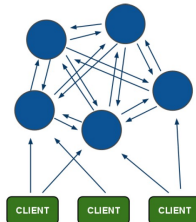
Key	Value
K1	AAA,BBB,CCC
K2	AAA,BBB
K3	AAA,DDD
K4	AAA,2,01/01/2015
K5	3,ZZZ,5623



# Keystore

## Uses

- Typical key value store
- High speed and in-memory storage suitable for caching applications
- Implementing application queues
- Can scale in a cluster



# Layout

## Εισαγωγή

Έννοιες

Relational

## Τύποι NoSQL βάσεων

Document

Keystore

**Column**

Graph

## Use Case

Wikipedia

Twitter

Others

## Conclusion

Conclusions

## Column

### Most popular

HBase, Cassandra, Hypertable

### Θετικά

- Horizontal Scaling
- Suitable for big data
- Usually build-in support for versioning and compression

## Column

### Most popular

HBase, Cassandra, Hypertable

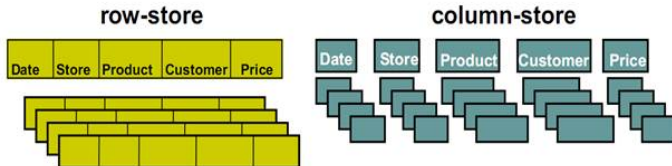
### Θετικά

- Horizontal Scaling
- Suitable for big data
- Usually build-in support for versioning and compression

### Αρνητικά

- schema based on how data is queried (know how data is to be used)
- Only suitable for large setups (minimum 5 nodes)

## Column based characteristics



## Column

- Data is stored at low level to be accessed efficiently as rows.
- Row based access needs to access much more data to get a column.
- Column-oriented storage makes sense for certain problems (eg range aggregates, analytic queries)
- Analytic queries can run orders of magnitudes quicker

# Layout

## Εισαγωγή

Έννοιες

Relational

## Τύποι NoSQL βάσεων

Document

Keystore

Column

Graph

## Use Case

Wikipedia

Twitter

Others

## Conclusion

Conclusions

# Graph

Most popular

Neo4J



# Graph

Most popular

Neo4J



## Θετικά

- OO designs
- Modeling interconnections
- Transactional
- Many language bindings and REST
- High Availability



# Graph

Most popular

Neo4J



## Θετικά

- OO designs
- Modeling interconnections
- Transactional
- Many language bindings and REST
- High Availability

## Αρνητικά

- Bad at partitioning (bad scale out)

# Neo4J INSERT

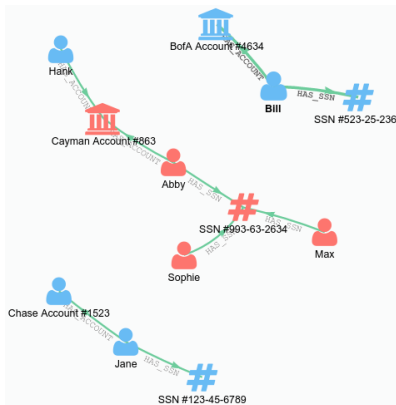
```
driver = GraphDatabase.driver("bolt://localhost", auth=basic_auth("neo4j", "test"))
session = driver.session()
```

```
# Insert data
```

```
insert_query = '''
CREATE (hank:Person {name:"Hank"}),
(abby:Person {name:"Abby"}),
(max:Person {name:"Max"}),
(sophie:Person {name: "Sophie"}),
(jane:Person {name: "Jane"}),
(bill:Person {name: "Bill"}),
(ssn993632634:SSN {number: 993632634}),
(ssn123456789:SSN {number: 123456789}),
(ssn523252364:SSN {number: 523252364}),
(chase:Account {bank: "Chase", number: 1523}),
(bofa:Account {bank: "Bank of America", number: 4634}),
(cayman:Account {bank: "Cayman", number: 863}),
(bill)-[:HAS_SSN]->(ssn523252364),
(bill)-[:HAS_ACCOUNT]->(bofa),
(jane)-[:HAS_SSN]->(ssn123456789),
(jane)-[:HAS_ACCOUNT]->(chase),
(hank)-[:HAS_ACCOUNT]->(cayman),
(abby)-[:HAS_ACCOUNT]->(cayman),
(abby)-[:HAS_SSN]->(ssn993632634),
(sophie)-[:HAS_SSN]->(ssn993632634),
(max)-[:HAS_SSN]->(ssn993632634)
'''
```

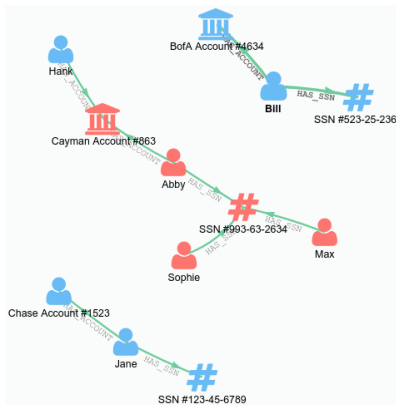
```
session.run(insert_query)
```

## Transitive Closure



Given suspicions about Hank, find related information to investigate.

## Transitive Closure



Given suspicions about Hank, find related information to investigate.

MATCH

$(n:Person) - [*] - (o)$

WHERE

$n.name = "Hank"$

RETURN

$o$

Hank shares an account with Abby. Abby shares a SSN with Sophie and Max. Given that we suspect Hank may be involved in fraudulent activity, we can flag the Cayman account, Sophie, and Abby as possible fraudulent entities.

## Neo4J Transitive Closure

```
query = '''  
MATCH (n:Person)-[*]-(o)  
WHERE n.name = {name}  
RETURN DISTINCT o AS other  
'''  
  
results = session.run(query, parameters={"name": "Hank"})  
for record in results:  
    print(record["other"])
```

## Neo4J Transitive Closure

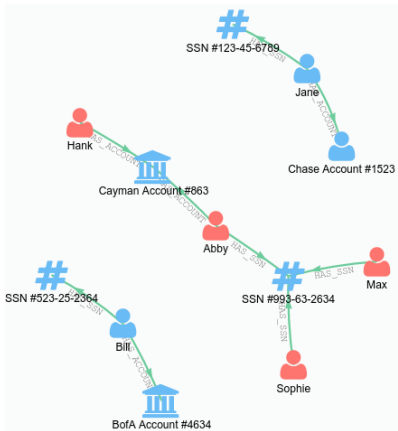
```
query = '''  
MATCH (n:Person)-[*]-(o)  
WHERE n.name = {name}  
RETURN DISTINCT o AS other  
'''
```

```
results = session.run(query, parameters={"name": "Hank"})  
for record in results:  
    print(record["other"])
```

output:

```
<Node id=95 labels=set([u'Account']) properties={u'number': 863, u'bank': u'Cayman'}>  
<Node id=85 labels=set([u'Person']) properties={u'name': u'Abby'}>  
<Node id=90 labels=set([u'SSN']) properties={u'number': 993632634}>  
<Node id=87 labels=set([u'Person']) properties={u'name': u'Sophie'}>  
<Node id=86 labels=set([u'Person']) properties={u'name': u'Max'}>
```

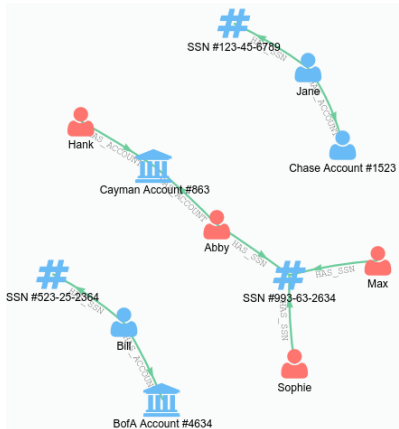
# Investigation Targeting



Fraud rings often share fraudulent identifying information. Any person with connections to more than two entities in the graph are suspicious. Find large cliques to investigate further.



# Investigation Targeting



Fraud rings often share fraudulent identifying information. Any person with connections to more than two entities in the graph are suspicious. Find large cliques to investigate further.

MATCH

$(n: \text{Person}) - [*] - (o)$

WITH

$n,$

$\text{count}(\text{DISTINCT } o) \text{ AS size}$

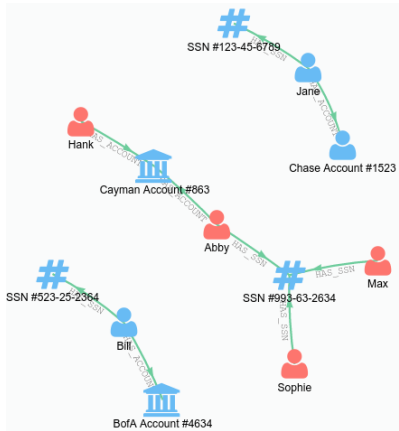
WHERE

$\text{size} > 2$

RETURN

$n$

# Investigation Targeting



Fraud rings often share fraudulent identifying information. Any person with connections to more than two entities in the graph are suspicious. Find large cliques to investigate further.

MATCH

$(n: \text{Person}) - [*] - (o)$

WITH

$n,$

$\text{count}(\text{DISTINCT } o) \text{ AS size}$

WHERE

$\text{size} > 2$

RETURN

$n$

Sophie, Max and Abby all share a SSN, which is suspicious. Hank is also suspicious because he is sharing an account with Abby.

## Neo4J Investigation Targeting

```
targeting_query = """  
MATCH (n:Person)-[*]-(o)  
WITH n, count(DISTINCT o) AS size  
WHERE size > 2  
RETURN n  
"""  
  
results = session.run(targeting_query)  
for record in results:  
    print(record["n"])
```

## Neo4J Investigation Targeting

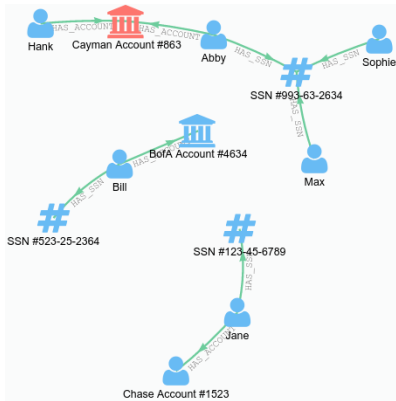
```
targeting_query = """
MATCH (n:Person)-[*]-(o)
WITH n, count(DISTINCT o) AS size
WHERE size > 2
RETURN n
"""
```

```
results = session.run(targeting_query)
for record in results:
    print(record["n"])
```

output:

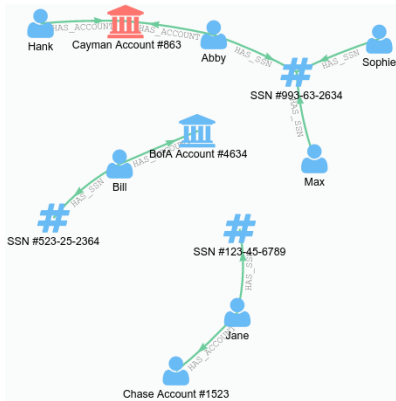
```
<Node id=86 labels=set([u'Person']) properties={u'name': u'Max'}>
<Node id=85 labels=set([u'Person']) properties={u'name': u'Abby'}>
<Node id=87 labels=set([u'Person']) properties={u'name': u'Sophie'}>
<Node id=84 labels=set([u'Person']) properties={u'name': u'Hank'}>
```

# Fast Insights



Given that we've identified SSN 993-63-2634 as suspicious, find all associated accounts.

# Fast Insights



Given that we've identified SSN 993-63-2634 as suspicious, find all associated accounts.

MATCH

```
(ssn : SSN) <- [:HAS_SSN] - (: Person) -  
[:HAS_ACCOUNT] -> (acct : Account)
```

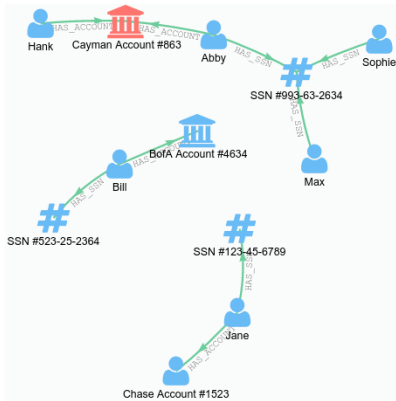
WHERE

```
ssn . number = 993632634
```

RETURN

```
acct
```

# Fast Insights



Given that we've identified SSN 993-63-2634 as suspicious, find all associated accounts.

MATCH

```
(ssn : SSN) <- [:HAS_SSN] - (: Person) -  
[:HAS_ACCOUNT] -> (acct : Account)
```

WHERE

```
ssn . number = 993632634
```

RETURN

```
acct
```

We see that the Cayman account #863 is the only account where a Person using this SSN owns the account.

## Neo4J Fast Insights

```
query = """  
MATCH (ssn:SSN)-[:HAS_SSN]-(:Person)-[:HAS_ACCOUNT]->(acct)  
WHERE ssn.number = {flagged_ssn}  
RETURN acct  
"""
```

```
results = session.run(query, parameters={"flagged_ssn": 999999999})  
for record in results:  
    print(record["acct"])
```



## Neo4J Fast Insights

```
query = """
MATCH (ssn:SSN)-[:HAS_SSN]-(:Person)-[:HAS_ACCOUNT]->(acct)
WHERE ssn.number = {flagged_ssn}
RETURN acct
"""
```

```
results = session.run(query, parameters={"flagged_ssn": 999})
for record in results:
    print(record["acct"])
```

output:

```
<Node id=95 labels=set([u'Account']) properties={u'number': 863, u'bank': u'Cayman'}>
```

Εισαγωγή  
○○○○○○○○○○  
○○○○○○○○○○○○

Τύποι NoSQL βάσεων  
○○○○○○○○○○○○  
○○○  
○○○  
○○○  
○○○○○○○○○

Use Case  
●○○○  
○○  
○○  
○○

Conclusion  
○○○○○

# Layout

## Εισαγωγή

Έννοιες

Relational

## Τύποι NoSQL βάσεων

Document

Keystore

Column

Graph

## Use Case

Wikipedia

Twitter

Others

## Conclusion

Conclusions

# Wikipedia

Μέγεθος

## Design Desisions

- 99-percent availability
- Geographic distribution
- Relaxed policy on security and data Transactions

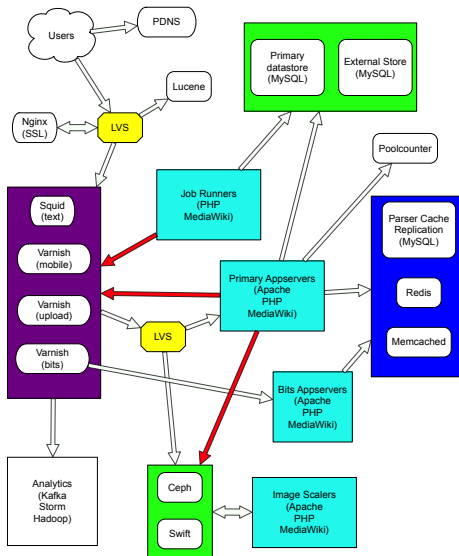
# Wikipedia

## Μέγεθος

### Design Desisions

- 99-percent availability
- Geographic distribution
- Relaxed policy on security and data Transactions
- 50,000 http requests per second
- 80,000 SQL queries per second
- 7 million registered users
- 18 million page objects in the English version
- 250 million page links
- 220 million revisions
- 1.5 terabytes of compressed data

# Wikipedia (Architecture)



**Internet frontend**

**Backend**

**Core services**

**Miscellaneous services**

**Other sites**

**mobile**

**donate**

**Main cluster (wikipedia, wiktionary, etc)**

**Florida**

**Amsterdam**

**PDF**

**Toolserver**

**Database dumps**

**Search**

**Apaches**

**Text storage**

**Databases**

**File servers**

**Image scalers**

**IRC**

**Subversion**

**Nagios**

**Ganglia**

**Mail**

**DNS**

**LDAP/IS**

**Logging**

**APT repositories**

**NTP, SSH**

**Network tools**

**Batch jobs**

**Scratch hosts**

**MySQL DB**

**Slave DB**

**Apache**

**Test equal**

**Upload equal**

**Load Balancer**

**Search host**

**Search crawler**

**File server**

**Image scaler**

**Web service**

**Mail/IRC service**

**Network snitch**

Εισαγωγή  
○○○○○○○○○○  
○○○○○○○○○○○○○○

Τύποι NoSQL βάσεων  
○○○○○○○○○○○○○○  
○○○  
○○○  
○○○  
○○○○○○○○○

Use Case  
○○○○  
●○  
○○

Conclusion  
○○○○○

# Layout

## Εισαγωγή

Έννοιες

Relational

## Τύποι NoSQL βάσεων

Document

Keystore

Column

Graph

## Use Case

Wikipedia

**Twitter**

Others

## Conclusion

Conclusions

# Twitter

Multiple databases



# Twitter

## Multiple databases

MySQL primary storage of Tweets and Users (custom fork)

# Twitter

## Multiple databases

- MySQL primary storage of Tweets and Users (custom fork)
- FlockDB in-house graph database, store social graph information (following, etc)

# Twitter

## Multiple databases

**MySQL** primary storage of Tweets and Users (custom fork)

**FlockDB** in-house graph database, store social graph information (following, etc)

**Memcached** "heavily modified" fork of Memcached 1.4.4

# Twitter

## Multiple databases

**MySQL** primary storage of Tweets and Users (custom fork)

**FlockDB** in-house graph database, store social graph information (following, etc)

**Memcached** "heavily modified" fork of Memcached 1.4.4

**Cassandra** services like Snowflake for quickly generating unique identifiers

# Twitter

## Multiple databases

**MySQL** primary storage of Tweets and Users (custom fork)

**FlockDB** in-house graph database, store social graph information (following, etc)

**Memcached** "heavily modified" fork of Memcached 1.4.4

**Cassandra** services like Snowflake for quickly generating unique identifiers

**Gizzard** in-house scala framework for building custom distributed databases (with arbitrary storage technology)

# Twitter

## Multiple databases

**MySQL** primary storage of Tweets and Users (custom fork)

**FlockDB** in-house graph database, store social graph information (following, etc)

**Memcached** "heavily modified" fork of Memcached 1.4.4

**Cassandra** services like Snowflake for quickly generating unique identifiers

**Gizzard** in-house scala framework for building custom distributed databases (with arbitrary storage technology)

**Apache Lucene** search index

# Twitter

## Multiple databases

**MySQL** primary storage of Tweets and Users (custom fork)

**FlockDB** in-house graph database, store social graph information (following, etc)

**Memcached** "heavily modified" fork of Memcached 1.4.4

**Cassandra** services like Snowflake for quickly generating unique identifiers

**Gizzard** in-house scala framework for building custom distributed databases (with arbitrary storage technology)

**Apache Lucene** search index

**HBase and Hadoop** analytics running over data every day

# Twitter

## Multiple databases

**MySQL** primary storage of Tweets and Users (custom fork)

**FlockDB** in-house graph database, store social graph information (following, etc)

**Memcached** "heavily modified" fork of Memcached 1.4.4

**Cassandra** services like Snowflake for quickly generating unique identifiers

**Gizzard** in-house scala framework for building custom distributed databases (with arbitrary storage technology)

**Apache Lucene** search index

**HBase and Hadoop** analytics running over data every day

**Redis** experimental timeline storage technology



# Layout

## Εισαγωγή

Έννοιες

Relational

## Τύποι NoSQL βάσεων

Document

Keystore

Column

Graph

## Use Case

Wikipedia

Twitter

**Others**

## Conclusion

Conclusions

# Facebook

## Size

- >200M active users
- >100M users log on to Facebook at least once each day
- >30M users update their statuses at least once each day

# Facebook

## Size

- >200M active users
- >100M users log on to Facebook at least once each day
- >30M users update their statuses at least once each day

## Databases

MySQL primary storage (4000 shards, 9000 memcached)

# Facebook

## Size

- >200M active users
- >100M users log on to Facebook at least once each day
- >30M users update their statuses at least once each day

## Databases

**MySQL** primary storage (4000 shards, 9000 memcached)

**Hive** simple summarization jobs, business intelligence, machine learning etc.

# Facebook

## Size

- >200M active users
- >100M users log on to Facebook at least once each day
- >30M users update their statuses at least once each day

## Databases

**MySQL** primary storage (4000 shards, 9000 memcached)

**Hive** simple summarization jobs, business intelligence, machine learning etc.

**Cassandra** Currently used for Facebook's private messaging.



Εισαγωγή	Τύποι NoSQL βάσεων	Use Case	Conclusion
οοοοοοοοοο	οοοοοοοοοοοο	οοοο	●οοοο
οοοοοοοοοοοοοο	οοο	οο	
	οοο	οο	
	οοοοοοοοοο		

# Layout

## Εισαγωγή

Έννοιες

Relational

## Τύποι NoSQL βάσεων

Document

Keystore

Column

Graph

## Use Case

Wikipedia

Twitter

Others

## Conclusion

Conclusions

## What to use

### Question:

Which **one** DB technology should I use?



## What to use

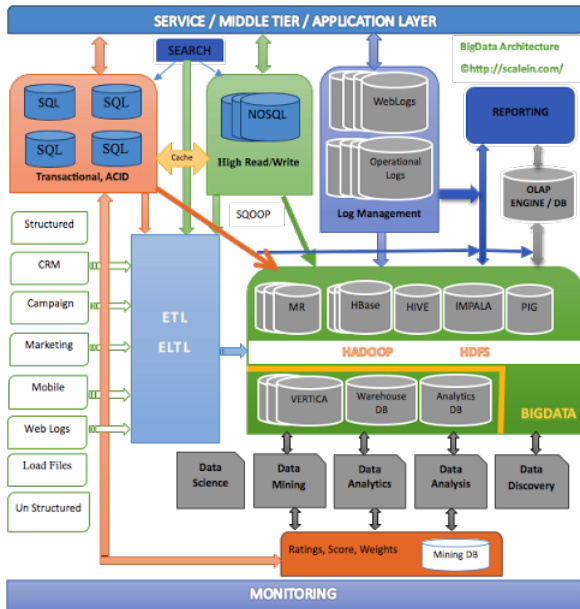
### Question:

Which **one** DB technology should I use?

### Answer:

Why use **one**? Depending on your scale you may need to use all of them!

# Typical BigData Architecture



## Conclusions

### Design requirements

- Query types (flexibility, relations, hierarchies)

## Conclusions

### Design requirements

- Query types (flexibility, relations, hierarchies)
- Data form (schema)

## Conclusions

### Design requirements

- Query types (flexibility, relations, hierarchies)
- Data form (schema)
- Data size (for sharding etc.)

## Conclusions

### Design requirements

- Query types (flexibility, relations, hierarchies)
- Data form (schema)
- Data size (for sharding etc.)
- Data inter-linkage

## Conclusions

### Design requirements

- Query types (flexibility, relations, hierarchies)
- Data form (schema)
- Data size (for sharding etc.)
- Data inter-linkage
- Web friendly

## Conclusions

### Design requirements

- Query types (flexibility, relations, hierarchies)
- Data form (schema)
- Data size (for sharding etc.)
- Data inter-linkage
- Web friendly
- Data analysis



## Conclusions

### Design requirements

- Query types (flexibility, relations, hierarchies)
- Data form (schema)
- Data size (for sharding etc.)
- Data inter-linkage
- Web friendly
- Data analysis
- Διαθεσιμότητα (9s)

## Conclusions

### Design requirements

- Query types (flexibility, relations, hierarchies)
- Data form (schema)
- Data size (for sharding etc.)
- Data inter-linkage
- Web friendly
- Data analysis
- Διαθεσιμότητα (9s)
- Requests per second (concurrent users)

## Conclusions

### Design requirements

- Query types (flexibility, relations, hierarchies)
- Data form (schema)
- Data size (for sharding etc.)
- Data inter-linkage
- Web friendly
- Data analysis
- Διαθεσιμότητα (9s)
- Requests per second (concurrent users)
- Read vs Write percentage

## Conclusions

### Design requirements

- Query types (flexibility, relations, hierarchies)
- Data form (schema)
- Data size (for sharding etc.)
- Data inter-linkage
- Web friendly
- Data analysis
- Διαθεσιμότητα (9s)
- Requests per second (concurrent users)
- Read vs Write percentage
- Policy on security, transactions (consistency, corruption)

## Conclusions

### Design requirements

- Query types (flexibility, relations, hierarchies)
- Data form (schema)
- Data size (for sharding etc.)
- Data inter-linkage
- Web friendly
- Data analysis
- Διαθεσιμότητα (9s)
- Requests per second (concurrent users)
- Read vs Write percentage
- Policy on security, transactions (consistency, corruption)
- Geographic distribution

Εισαγωγή

oooooooooooo

Τύποι NoSQL βάσεων

oooooooooooo  
ooo  
ooo  
oooooooooooo

Use Case

oooo  
oo  
oo

Conclusion

oooo●

*Thank you!*